

OBJECT CLASSIFICATION USING CNN -Acc 0.8235

Given a set of images that are all labeled with a single category, I predicted these categories for a set of images and measure the accuracy of the predictions and generated the confusion matrix.

Instead of trying to specify what every one of the image categories of interest look like directly in code, they provide the computer with many examples of each image class and then develop learning algorithms that look at these examples and learn about the visual appearance of each class. In other words, they first accumulate a training dataset of labeled images, then feed it to the computer in order for it to get familiar with the data.

Given that fact, the complete image classification pipeline can be formalized as follows:

- Our input is a training dataset that consists of 421 images, each labeled with one of 4 different classes.
- Then, we use this training set to train a classifier to learn what every one of the classes looks like.
- In the end, we evaluate the quality of the classifier by asking it to predict labels for a new set of images that it has never seen before. We will then compare the true labels of these images to the ones predicted by the classifier.
- Finally calculated the confusion matrix.

Here I used a 5 layers Sequential Convolutional Neural Network for object classification trained on the dataset downloaded from IKEA site (<https://www.ikea.com/>). I chose to build it with keras API (Tensorflow backend) which is very intuitive. Firstly, I will prepare the data then I will focus on the CNN modeling and evaluation.

I achieved **82.35%** of accuracy with this CNN trained.

Define the model

I used the Keras Sequential API, where you have just to add one layer at a time, starting from the input.

Here I used 5 convolution layers. The first is the convolutional (Conv2D) layer. It is like a set of learnable filters. I chose to set 32 filters for the first three conv2D layers and 64 filters for the next ones. The last conv2D layer have 128 filters. Each filter transforms a part of the image

(defined by the kernel size) using the kernel filter. The kernel filter matrix is applied on the whole image. Filters can be seen as a transformation of the image.

The second important layer in CNN is the pooling (MaxPool2D) layer. This layer simply acts as a downsampling filter. It looks at the 2 neighboring pixels and picks the maximal value. These are used to reduce computational cost, and to some extent also reduce overfitting. We have to choose the pooling size (i.e the area size pooled each time) more the pooling dimension is high, more the downsampling is important.

Combining convolutional and pooling layers, CNN are able to combine local features and learn more global features of the image.

Dropout is a regularization method, where a proportion of nodes in the layer are randomly ignored (setting their weights to zero) for each training sample. This technique also improves generalization and reduces the overfitting.

'Relu' is the rectifier (activation function $\max(0, x)$). The rectifier activation function is used to add non linearity to the network.

The Flatten layer is used to convert the final feature maps into a one single 1D vector. This flattening step is needed so that you can make use of fully connected layers after some convolutional/maxpool layers. It combines all the found local features of the previous convolutional layers. In the end I used the features in two fully-connected (Dense) layers which is just artificial neural networks (ANN) classifier. In the last layer(Dense(10,activation="softmax")) the net outputs distribution of probability of each class.

Once our layers are added to the model, we need to set up a score function, a loss function and an optimisation algorithm. I used a specific form for categorical classifications (>2 classes) called the "categorical_crossentropy" as a loss function.

I choosed Adam as optimizer, it is a Straightforward to implement, Computationally efficient, Little memory requirements and Well suited for problems that are large in terms of data and/or parameters.

Comparision with other architectures :-

I also created another two architectures for the object classification. They given very poor accuracy compared with this architecture(explained above). One of it given 68.235 % and another one given 24.7 % accuracy. They also specified in the folder **Another_Architectures** with the models.

Observations

The parameters of the model:-

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 398, 398, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 199, 199, 32)	0
dropout_1 (Dropout)	(None, 199, 199, 32)	0
conv2d_2 (Conv2D)	(None, 197, 197, 32)	9248
max_pooling2d_2 (MaxPooling2D)	(None, 98, 98, 32)	0
dropout_2 (Dropout)	(None, 98, 98, 32)	0
conv2d_3 (Conv2D)	(None, 96, 96, 32)	9248
max_pooling2d_3 (MaxPooling2D)	(None, 48, 48, 32)	0
dropout_3 (Dropout)	(None, 48, 48, 32)	0
conv2d_4 (Conv2D)	(None, 46, 46, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 23, 23, 64)	0
dropout_4 (Dropout)	(None, 23, 23, 64)	0
conv2d_5 (Conv2D)	(None, 21, 21, 128)	73856
dropout_5 (Dropout)	(None, 21, 21, 128)	0
flatten_1 (Flatten)	(None, 56448)	0
dense_1 (Dense)	(None, 128)	7225472
dropout_6 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 4)	516
Total params: 7,337,156		
Trainable params: 7,337,156		
Non-trainable params: 0		

The total parameters of the model = 73,37,156

Accuracy = 82.35 %

Confusion Matrix:-

Confusion Matrix :

```
[[51  0  0  0]
 [ 0 28  0  0]
 [ 0  0 62  0]
 [ 0  0  0 29]]
```

Accuracy Score : 1.0

Report :

	precision	recall	f1-score	support
0	1.00	1.00	1.00	51
1	1.00	1.00	1.00	28
2	1.00	1.00	1.00	62
3	1.00	1.00	1.00	29
micro avg	1.00	1.00	1.00	170
macro avg	1.00	1.00	1.00	170
weighted avg	1.00	1.00	1.00	170